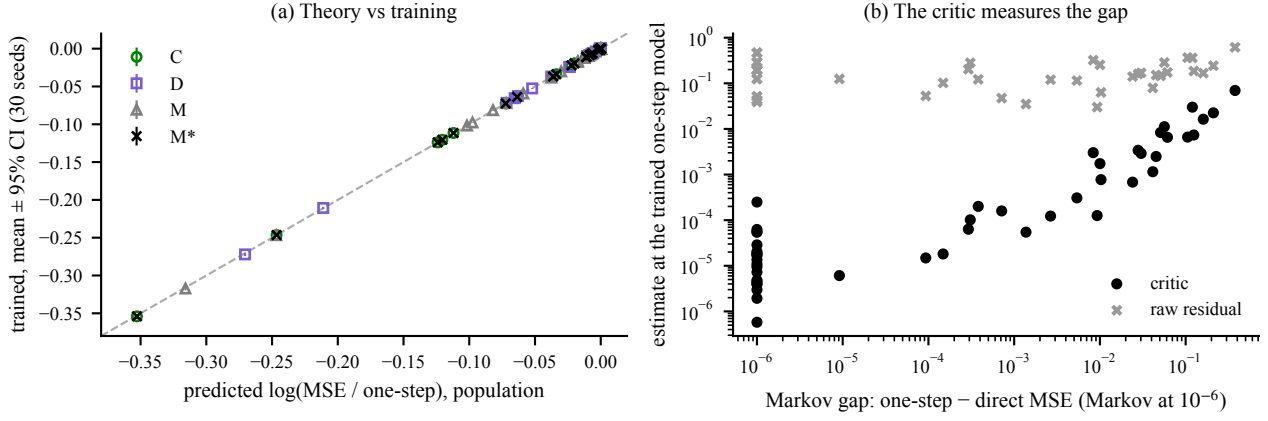
(a) Theory vs training
trained, mean ± 95% CI (30 seeds)
predicted log(MSE / one-step), population
C
D
M
M*
0.00
−0.05
−0.10
−0.15
−0.20
−0.25
−0.30
−0.35
(b) The critic measures the gap
estimate at the trained one-step model
Markov gap: one-step − direct MSE (Markov at 10^{-6})
critic
raw residual
10^{0}
10^{-1}
10^{-2}
10^{-3}
10^{-4}
10^{-5}
10^{-6}